

Systemic Compliance Platform: Where We're Headed, and Why It Matters

A plain-language summary of the architecture assessment: business goals, target design, key risks, and recommended next steps.

Prepared for: Systemic Compliance Leadership

1 Executive Brief

Systemic Compliance is really building **two things** today: a platform customers use (**SC.ORB**), and an AI-driven engine that produces and updates compliance content behind the scenes (**SC.IMS**). Right now, these are not one connected product; the engine lives largely inside a single AI vendor's environment, separate from the customer-facing platform.

The recommended direction is to make these **feel like one product to customers**, while keeping the intelligence engine as an **independently owned and operated** capability. This is what turns the company's compliance expertise into a real, defensible, sellable business asset instead of something borrowed from a vendor's toolset.

The assessment also surfaced two important risk areas that need attention before scaling further: **heavy reliance on a single vendor** for how the platform is built and run, and **no independent way yet to verify** that AI-generated content is accurate before it reaches a customer. Neither blocks the business today, but both need a plan.

STRATEGIC DECISION

One product, two independently owned engines

BIGGEST BUSINESS RISK

Single-vendor dependency (Lovable)

BIGGEST TRUST GAP

No independent AI quality check yet

RECOMMENDED FIRST MOVE

Harden the foundation before adding new capability

What We're Building Toward

In plain terms, here is what the future-state platform is meant to deliver for the business and its customers.

1

One Trusted Platform

Customers experience a single, coherent compliance product, not two disconnected systems stitched together.

2

AI That Moves Faster Than Manual Work

Regulatory changes and compliance content are drafted and routed automatically, instead of being written from scratch every time.

3

Provable Accuracy

Every AI-produced answer is checked by an independent process before it reaches a customer, so trust is earned, not assumed.

4

Room to Grow

New compliance areas (training, incident tracking, contractor management, safety programs) can be added without rebuilding the foundation.

5

A Platform the Company Owns

The core intelligence, the thing that actually creates value, is fully owned and controlled, not dependent on a single outside vendor.


6

Ready for Investment or Sale

Whether the company keeps building for years or explores a sale or partnership, the platform should look and run like an independent, well-managed asset.

3 The Future-State Picture

How information should flow once the target design is in place: from regulatory source, through the AI engine, past an independent quality check, into the platform customers use.



Risks to Address

Risks are grouped into business risks (impact on the company and its options) and technical risks (impact on how safely and quickly the platform can change).

Business Risks

HIGH**Single-Vendor Dependency**

A large part of how the platform is built, deployed, and run today relies on one outside provider. This could limit flexibility, slow negotiations, or reduce the platform's value in a sale.

HIGH**No Independent Check on AI Output**

There is currently no automated, independent system that verifies AI-generated content before it's used. This is a trust and liability exposure as usage grows.

MEDIUM**Confidence Relies on Manual Review**

Most confidence in the system today comes from people testing it by hand rather than automated checks. This is slower and costlier as the platform scales.

MEDIUM**Harder to Prove Ownership for a Deal**

If a future sale, partnership, or investment conversation happens, it will take extra work today to show the platform is fully independent and company-controlled.

Technical Risks

HIGH**Thin Automated Testing**

Only a small amount of automated "click-through" testing exists. There is no broader automated safety net, so changes carry more risk than they should.

MEDIUM**Complex Logic Buried in the Database**

A large amount of business logic (nearly 200 automated rules and close to 300 triggered actions) lives directly inside the database. It works, but it's hard to inspect and change safely.

MEDIUM**No Automated Grading for AI Content**

There isn't yet a repeatable, automatic way to measure whether AI-generated content is accurate and consistent before people see it.

MEDIUM**No Support for Long-Running AI Work**

The current hosting setup can't run long AI tasks directly. A separate, purpose-built engine is required to support serious AI workloads reliably.

Recommended Next Steps

A phased approach that builds trust and independence first, then expands capability, rather than adding new features on top of an unproven foundation.

● Establish a Foundation of Trust

Put basic automated testing and a clear inventory of what needs to change in place, so future work is safer and easier to measure.

● Stand Up the Independent AI Engine

Build the compliance intelligence engine as its own, company-owned system, not tied to a single vendor's build tools.

● Create the Quality Gate

Introduce an independent way to check AI-generated content before it's trusted and shown to customers.

● Reduce Vendor Dependency

Systematically replace vendor-tied pieces (build and deploy, sign-in, email, and AI routing) with company-owned alternatives, one piece at a time.

● Expand Platform Capabilities

Once the foundation is solid, confidently grow into new compliance areas: training, incident management, contractor programs.

● Keep a Business-Readiness Package Current

Maintain clear documentation of ownership, architecture, and operations, so the company is always ready for investment, partnership, or sale conversations.