

Call Prep — Matt (Mehmet), 2026-08-28

Purpose of this call (as you frame it going in): understand where Matt is actually taking the SC.ORB build over the next few months — not just what's shipped — so you can sanity-check the roadmap against the capability/positioning work in `2026-08-27-orb-addon-capability-assessment-and-positioning.md` before Kevin/Robin lock priorities. Per the 2026-07-08 whiteboard session, this is the "technical deep-dive with Matt to assess architecture and gaps" that was flagged as an action item — if Kevin/Robin aren't on this call, keep that in mind for how much you commit to on SC's behalf versus just gathering information.

0. Background context (not today's agenda) — Robin (Aug 24) and Matt (Aug 25), "Scoping the next phase"

This landed in email right before today's call. It's background awareness, not something to work through on this call — Robin's asks are explicitly for a separate working session next week. Today is still just the exploratory conversation about where Matt is taking the platform.

Robin's Aug 24 email (to you, cc Kevin + Matt) moves the engagement from direction to execution — they got the first invoice and want a working session, framed hard around cost-as-a-design-constraint (they're a startup, this is a real stretch for their budget). What she's asking Improving to bring to that session:

1. A ranked recommendation — of the six roadmap steps in the Platform Strategy Brief, which 2-3 give the most risk reduction / value protection per dollar, rest deferred.
2. Scope + cost per item, split explicitly between **Improving hours vs. Matt's hours** — she wants your hours concentrated on what's genuinely differentiated (AI quality/eval work specifically) and as much build as possible done in-house by Matt.
3. A **scoped assessment of "Reduce Vendor Dependency"** specifically — she flagged it as reading like the largest item on the list and wants sizing, a lighter-weight alternative if one exists, and your honest read on what a buyer/investor would actually discount for if left as-is.
4. Direction broken into defined, executable tasks for Matt — not just "what should happen."
5. A defined working model for how Matt gets direction between sessions without billable overhead every time.

This is not today's call — she asked for a separate working session next week, Kevin permitting. It's just the lens Matt is operating under right now.

Matt's Aug 25 reply (same thread, to you, cc Kevin + Robin) is his technical pushback on the brief, and it's worth reading as pushback, not just input:

SC.IMS is not a separate product — he's explicit: it's internal tooling (consulting/BD team use), named after Kevin's early Claude AI work, not a customer-facing platform. Any framing of SC.IMS as a distinct thing sitting alongside SC.ORB should be corrected — it's positioned as **feeding into** ORB, not competing with or separate from it.

SC.IMS-to-ORB integration is already in motion, not theoretical: they're actively wiring SC.IMS tooling into a customizable MOC process inside ORB, with a **customer-demo version expected in ~2 weeks from Aug 25 (so roughly Sept 8)**.

He directly disputes the brief's "Reduce Vendor Dependency" conclusion on Lovable. His position: Lovable is a dev tool, doesn't own their code/data/infra/ability to operate the platform, and he's asking you point-blank to explain the specific dependency, its impact, and the basis for that conclusion. He's also asserting SC.ORB was designed from the start as independently owned/portable.

He's offering to take technical questions directly (phone/Teams) rather than routing everything through email.

What this means for today's call: none of Robin's ask items are today's business — that's the separate working session next week, and it's not on you to resolve or defend anything from the brief in this conversation. The one thing worth carrying in is a corrected mental model: treat SC.IMS as ORB's internal AI tooling lineage, not a parallel system, when you ask the MOC/SC-LINK questions in section 2 below. Today is purely about understanding where Matt is actually taking the platform.

1. Where is the roadmap actually headed? (the main thing you need out of this call)

His own milestone doc (`mem/roadmap/scorb-milestones.md`) bundles **v8.0 "Compliance Expansion"** as Dashboards (M7) + SC-Score/Benchmarking (M8) + Audits (M9) + LMS (M10) + AI layer expansion (M11) — four genuinely large builds with only M7 in progress and M9 explicitly "NOT STARTED." You need his actual sequencing, not the milestone doc's bundling order, because the bundling order and the business priority order may not match.

Ask directly: in what order are M8/M9/M10/M11 actually getting built, and why that order? (His answer tells you whether engineering sequencing or business value is driving the roadmap.)

Push specifically on **LMS (M10)** — per the capability assessment, this is the single highest-leverage gap: Competence Architect already generates real, production content (SCORM, bilingual, a 62-course library) with nowhere real to land, because the LMS module today has **zero domain database tables** (`docs/lms/LIMITATIONS.md`: hard-coded fake learner, no real ZIP extraction, client-trusted scores). Ask what "done" looks like for him here and whether he's scoped ingesting Competence Architect's actual output or is still thinking of LMS as a generic player.

Push on **SC-Score / Benchmarking (M8)** — SC.ORB already has a "Coming Soon" placeholder for this with nothing behind it. This is the most commercially visible gap (it's the centerpiece of the Compliance Intelligence & Proof theme). Ask what data model he's imagining and whether he's aware the scoring math itself already exists on the SC.IMS side (Conformance + Coherence axes, feeding an SCI) — you don't want him reinventing scoring logic that's already been solved.

Ask about **Audits (M9)** honestly — it's billable in the module registry today but has no schema at all. Does it get built this cycle, or is it aspirational flag-setting?

2. The MOC/SC-LINK question — handle carefully, this is the IP conversation

ORB's MOC module today is the change-request/CIA shell only — the actual ripple-prediction engine (SC-LINK) lives outside ORB as a standalone tool and is logged as protected trade-secret IP. This is worth surfacing now, not after it's half-wired in.

Update per Matt's Aug 25 email: SC.IMS-related tooling is already being integrated into ORB specifically for a customizable MOC process, with a customer-demo version targeted for ~2 weeks out from Aug 25 (roughly Sept 8). This is further along than "half-wired in" — worth asking directly whether this demo-targeted MOC integration includes SC-LINK's ripple/impact-prediction mechanism itself, or is scoped to the change-request/CIA shell plus SC.IMS-side onboarding/requirements-loading only (his email describes the latter — "supports client onboarding and loading client requirements into the platform" — which reads like it may not yet touch the actual ripple-prediction IP). Don't assume the trade-secret exposure question is resolved just because integration work has started; if anything, an active integration effort makes the API-surface/exposure question in the bullet below more urgent, not less.

Ask what his plan is for wiring MOC's full ripple/impact-prediction into ORB, and specifically how the mechanism will be exposed (or not) at the product surface — API responses, client-side code, exportable data. This matters because trade-secret protection depends on it never leaking through the product itself, not just on nobody talking about it.

Use this moment to surface, gently, whether Matt (and any contractor/dev he brings in) is under a **signed IP-assignment and confidentiality agreement** covering this specific asset. This was flagged as an open governance gap in the July 8 whiteboard session (in that case about the training corpus) — worth checking whether it's been closed since, because it directly affects whether SC-LINK's trade-secret status is actually defensible.

Don't need to resolve this on the call — just confirm it's on someone's list (probably counsel's, not Matt's) and isn't assumed solved.

3. Billing / payments — the blocker behind every new-segment idea

`subscription\_plan\_template` is a fully-modeled rate-card/entitlement schema with **no payment processor wired at all** — no Stripe, no invoice tables, no webhook handling. Several of the newer product ideas (Transaction Diligence, any one-time-fee SKU) are blocked on this specifically, not on missing regulatory content.

Ask if/when a payment processor gets wired, and whether he's planned for a **one-time-fee primitive** or only ever designed for recurring subscriptions. If it's subscription-only by design, that's a real constraint worth knowing about before anyone pitches a one-time diligence product.

4. AI layer — check the guardrail is still intentional, not just unaddressed

The AI layer today is nav-help/chat only via the Lovable AI Gateway, and `mem/index.md` records a **locked rule**: Lovable AI Gateway is the ***only*** AI provider, deep domain assistance (gap analysis, MOC/OQ assistants) is explicitly deferred to "Phase 7+."

Ask if that's still the constraint, or if it's loosening. Worth noting out loud (diplomatically) that the **SC.IMS side runs on Claude** for exactly the kind of domain reasoning (gap analysis, coherence scoring) ORB's AI layer is deferring — if ORB ever wants to expose any of that reasoning to a customer, the two systems are on different model providers today, and that's worth being an explicit decision rather than an accident of who built what first.

Correction per Matt's Aug 25 email: SC.IMS is not a separate system/product in his framing — it's the name for a set of internal tools (consulting/BD team use) that grew out of Kevin's early work with Claude AI, feeding into ORB rather than sitting alongside it as a parallel platform. Frame the model-provider question accordingly: it's "ORB's customer-facing AI layer is locked to Lovable while ORB's internal tooling runs on Claude," not "two separate products on two separate providers." The underlying point (different providers for different reasoning tasks, worth an explicit decision) still stands.

FYI, not for today — Matt is actively disputing the Vendor Dependency conclusion on Lovable from the Platform Strategy Brief in the same email: his position is Lovable is a dev tool only, doesn't own their code/data/infra, and he wants you to explain the specific dependency/impact/basis for calling it the platform's biggest risk. He's also asserting SC.ORB was architected from day one to be independently owned and portable. That's Robin's working-session territory next week, not this call — no need to engage it here, just be aware it's live.

5. Loose ends from the July 9 demo — quick status check

Did the **technical build documentation and security/infosec documentation** he committed to share ever get delivered? If not, ask again directly — you flagged real data-retention/audit-liability exposure (operator email domain hosting) and PCI/Stripe scope questions in that session that are still worth closing out.

Hosting/deployment footprint — as-built docs found no direct evidence of the claimed Azure hosting; the README says "host-neutral." Worth a quick, low-friction "where does this actually run today" so you have a real answer, not an inference.

Gemini vs. Lovable AI Gateway — the July demo said Gemini; the current codebase shows Lovable AI Gateway as the **only** allowed provider per the locked rule. Worth clarifying whether that changed deliberately or the two data points are from different points in time.

6. Known deferred/open items — worth a fast pass, not a deep dive

Quick list to run down verbally and see which ones have moved:

- Equivalency-granting workflow (OQ) — matrix exists, granting logic deferred.
- Cross-module MOC change-log writer — deferred.
- Auth email canonicalization collision bug — deferred.
- MFA recovery / host super-admin reset — deferred.
- Dashboard widget registry — landing incrementally per module; ask if this blocks or is independent of M8's scoring dashboard.

Quick-reference: what's actually shipped vs. stubbed (so you don't need the other doc mid-call)

Area	State
OQ, SWP, D&A, COI, ORB requirements repo	Shipped, billable
MOC	Change-request shell shipped; full cross-standard ripple not wired
SC.ORBard / Field Scorecard Verification	Shipped (QR scan, 3-state rollup)
LMS	Nav + player scaffolding only, **zero domain tables**
Audits	Nav entry + billable flag only, **zero schema**
SC-Score / Benchmarking	"Coming Soon" placeholder only
AI layer	Nav help/chat only, Lovable AI Gateway locked-in

Billing	Rate-card schema exists, **no payment processor wired**
SSO/SCIM/MFA	Shipped

What you want to walk away with

1. His real sequencing for M8–M11, in his own words, and why.
2. A concrete answer on LMS: is he building toward ingesting Competence Architect's output, or building a generic player that'll need a second integration pass later?
3. Confirmation the MOC/SC-LINK IP question is on someone's radar (even if not his to solve) — and whether the ~Sept 8 MOC/SC.IMS demo he committed to Robin/Kevin touches the actual ripple-prediction mechanism or stays scoped to onboarding/requirements-loading.
4. Whether payments/one-time-fee billing is anywhere on his roadmap.
5. Status of the outstanding docs (technical build + security/infosec) from July 9.